

Deep Learning Assignment-01

MTech (CS), IIIT Bhubaneswar
March - 2025



Student ID: A124010 Student Name: Shrikant Singh

Solutions of Deep Learning Assignment 01

1. : For a D -dimensional input vector, show that the optimal weights can be represented by the expression: 1

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$$

What is the possible estimation of \mathbf{w} ?

Solution:

To derive the optimal weights \mathbf{w} for a linear regression problem, we start with the least squares objective. Given a dataset with N samples, where \mathbf{X} is the $N \times D$ design matrix (each row corresponds to a D -dimensional input vector), \mathbf{t} is the $N \times 1$ target vector, and \mathbf{w} is the $D \times 1$ weight vector, the goal is to minimize the sum of squared errors:

$$E(\mathbf{w}) = \|\mathbf{t} - \mathbf{X}\mathbf{w}\|^2$$

Expanding the squared error term:

$$E(\mathbf{w}) = (\mathbf{t} - \mathbf{X}\mathbf{w})^T (\mathbf{t} - \mathbf{X}\mathbf{w})$$

Taking the derivative of $E(\mathbf{w})$ with respect to \mathbf{w} and setting it to zero for minimization:

$$\frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{t} - \mathbf{X}\mathbf{w}) = 0$$

Rearranging the equation:

$$\mathbf{X}^T \mathbf{t} - \mathbf{X}^T \mathbf{X} \mathbf{w} = 0$$

Solving for \mathbf{w} :

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{t}$$

Assuming $\mathbf{X}^T \mathbf{X}$ is invertible, the optimal weight vector \mathbf{w} is:

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$$

This is the least squares solution for the weight vector \mathbf{w} .

Estimation of \mathbf{w}

The expression $\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$ provides the optimal weights that minimize the sum of squared errors between the predicted values $\mathbf{X}\mathbf{w}$ and the target values \mathbf{t} . This is the best linear unbiased estimator (BLUE) under the assumptions of linear regression (e.g., no multicollinearity, homoscedasticity, and normally distributed errors).

Thus, the estimation of \mathbf{w} is given by:

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$$

2. : OR Gate Implementation Using a Single-Layer Neural Network

Solution:

OR Gate Using a Perceptron

The perceptron implements the OR logic through:

$$y = f(w_1 x_1 + w_2 x_2 + b)$$

Where:

- Initial weights: $w_1 = 1.5, w_2 = 2$
- Initial bias: $b = -2$
- Learning rate: $\eta = 0.5$
- Step activation:

$$f(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

OR Gate Truth Table

x_1	x_2	t
0	0	0
0	1	1
1	0	1
1	1	1

Learning Process

Epoch 1 - Initial Verification

1. Input (0,0):

$$z = 1.5(0) + 2(0) - 2 = -2 \Rightarrow \mathbf{y=0} \quad \checkmark$$

2. Input (0,1):

$$z = 1.5(0) + 2(1) - 2 = 0 \Rightarrow \mathbf{y=1} \quad \checkmark$$

3. Input (1,0):

$$z = 1.5(1) + 2(0) - 2 = -0.5 \Rightarrow \mathbf{y=0} \quad \times \text{ (Expected 1)}$$

Weight Update:

$$\Delta w_1 = \eta(t - y)x_1 = 0.5(1 - 0)(1) = 0.5$$

$$\Delta w_2 = \eta(t - y)x_2 = 0.5(1 - 0)(0) = 0$$

$$\Delta b = \eta(t - y) = 0.5(1 - 0) = 0.5$$

$$w_1 \leftarrow 1.5 + 0.5 = 2.0$$

$$w_2 \leftarrow 2 + 0 = 2.0$$

$$b \leftarrow -2 + 0.5 = -1.5$$

4. Input (1,1) with updated parameters:

$$z = 2(1) + 2(1) - 1.5 = 2.5 \Rightarrow \mathbf{y=1} \quad \checkmark$$

Epoch 2 - Verification with Updated Parameters

New Parameters: $w_1 = 2.0$, $w_2 = 2.0$, $b = -1.5$

1. Input (0,0):

$$z = 2(0) + 2(0) - 1.5 = -1.5 \Rightarrow \mathbf{y=0} \quad \checkmark$$

2. Input (0,1):

$$z = 2(0) + 2(1) - 1.5 = 0.5 \Rightarrow \mathbf{y=1} \quad \checkmark$$

3. Input (1,0):

$$z = 2(1) + 2(0) - 1.5 = 0.5 \Rightarrow \mathbf{y=1} \quad \checkmark$$

4. Input (1,1):

$$z = 2(1) + 2(1) - 1.5 = 2.5 \Rightarrow \mathbf{y=1} \quad \checkmark$$

Convergence Achieved

After weight adjustment in Epoch 1, the perceptron correctly classifies all OR gate inputs. The final parameters are:

$$\boxed{w_1 = 2.0, \quad w_2 = 2.0, \quad b = -1.5}$$

The decision boundary equation becomes:

$$2.0x_1 + 2.0x_2 - 1.5 = 0$$

3. Design a Perceptron algorithm to classify Iris flowers using either sepal or petal features and create a decision boundary.

Solution:

Perceptron Algorithm for Iris Classification

Objective: Classify two species of Iris flowers using:

- **Feature choice:** Sepal (length/width) *or* Petal (length/width)
- **Class pairs:** Setosa-Virginica, Setosa-Versicolor, or Versicolor-Virginica

Algorithm Implementation

1. Data Preparation

- Feature selection:

$$\text{Features} = \begin{cases} (\text{sepal length, sepal width}) & \text{if sepal mode} \\ (\text{petal length, petal width}) & \text{if petal mode} \end{cases}$$

- Binary encoding (example for Setosa vs Virginica):

$$y = \begin{cases} 1 & \text{Virginica} \\ 0 & \text{Setosa} \end{cases}$$

- Min-max normalization:

$$x_{\text{norm}} = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

2. Model Initialization

- Weights: $w_1, w_2 \sim \mathcal{U}(-0.01, 0.01)$
- Bias: $b = 0$
- Learning rate: $\eta = 0.1$ (default)

3. Training Phase

 For each epoch until convergence:

- (a) Compute activation:

$$z = w_1 x_1 + w_2 x_2 + b$$

- (b) Apply step function:

$$\hat{y} = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{otherwise} \end{cases}$$

- (c) Update parameters for misclassified samples:

$$\Delta w_i = \eta(y - \hat{y})x_i \quad (\text{for } i = 1, 2)$$

$$\Delta b = \eta(y - \hat{y})$$

4. **Decision Boundary** The separating line equation:

$$w_1x_1 + w_2x_2 + b = 0$$

Slope-intercept form:

$$x_2 = -\frac{w_1}{w_2}x_1 - \frac{b}{w_2}$$

Python Implementation

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

def get_user_choice():
    """Get user input for feature selection and class pair"""
    print("Available options:")
    print("1. Sepal features (length & width)")
    print("2. Petal features (length & width)")
    feature_choice = int(input("Enter feature choice (1 or 2): "))
    )

    print("\nClass pairs:")
    print("1. Setosa (0) vs Versicolor (1)")
    print("2. Setosa (0) vs Virginica (2)")
    print("3. Versicolor (1) vs Virginica (2)")
    class_pair = int(input("Enter class pair (1-3): "))

    return feature_choice, class_pair

def prepare_data(feature_choice, class_pair):
    """Load and prepare data based on user choices"""
    iris = load_iris()

    # Feature selection
    if feature_choice == 1:
        X = iris.data[:, :2] # Sepal features
        feature_names = ["Sepal Length", "Sepal Width"]
    else:
        X = iris.data[:, 2:] # Petal features
        feature_names = ["Petal Length", "Petal Width"]

    # Class pair selection
    class_mapping = {1: (0, 1), 2: (0, 2), 3: (1, 2)}
    class_a, class_b = class_mapping[class_pair]

    # Filter selected classes
    mask = np.logical_or(iris.target == class_a, iris.target ==
        class_b)
    X = X[mask]
    y = iris.target[mask]
```

```

# Convert to binary labels
y = np.where(y == class_b, 1, 0)

return X, y, feature_names, iris.target_names[class_a], iris.
        target_names[class_b]

def perceptron_train(X_train, y_train, lr=0.1, epochs=1000):
    """Train perceptron model"""
    weights = np.zeros(X_train.shape[1])
    bias = 0

    for epoch in range(epochs):
        for i in range(len(X_train)):
            z = np.dot(X_train[i], weights) + bias
            y_pred = 1 if z > 0 else 0
            error = y_train[i] - y_pred
            weights += lr * error * X_train[i]
            bias += lr * error

    return weights, bias

def evaluate_model(X_test, y_test, weights, bias):
    """Evaluate model on test data"""
    correct = 0
    for i in range(len(X_test)):
        z = np.dot(X_test[i], weights) + bias
        y_pred = 1 if z > 0 else 0
        if y_pred == y_test[i]:
            correct += 1
    return correct / len(X_test)

def plot_results(X_train, X_test, y_train, y_test, weights, bias,
                 feature_names, class_names):
    """Plot decision boundary with train/test points"""
    plt.figure(figsize=(10, 6))

    # Create mesh grid
    x_min, x_max = X_train[:, 0].min()-1, X_train[:, 0].max()+1
    y_min, y_max = X_train[:, 1].min()-1, X_train[:, 1].max()+1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100),
                          np.linspace(y_min, y_max, 100))

    # Calculate decision boundary
    Z = (weights[0]*xx + weights[1]*yy + bias > 0).astype(int)

    # Plot decision regions
    plt.contourf(xx, yy, Z, alpha=0.2, cmap='RdBu')

    # Plot training points
    plt.scatter(X_train[:, 0], X_train[:, 1], c=y_train,
                cmap='RdBu', edgecolors='k', label='Train', alpha
                =0.7)

    # Plot test points
    plt.scatter(X_test[:, 0], X_test[:, 1], c=y_test,

```

```

        cmap='RdBu', edgecolors='k', marker='x', s=100,
        label='Test')

plt.xlabel(f"{feature_names[0]} (standardized)")
plt.ylabel(f"{feature_names[1]} (standardized)")
plt.title(f"Perceptron: {class_names[0]} vs {class_names[1]}")
    )
plt.legend()
plt.show()

def main():
    # Get user choices
    feature_choice, class_pair = get_user_choice()

    # Prepare data
    X, y, feature_names, class_a, class_b = prepare_data(
        feature_choice, class_pair)

    # Split data (80-20)
    X_train, X_test, y_train, y_test = train_test_split(X, y,
        test_size=0.2, random_state=42)

    # Standardize features
    scaler = StandardScaler()
    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)

    # Train model
    weights, bias = perceptron_train(X_train, y_train)

    # Evaluate
    accuracy = evaluate_model(X_test, y_test, weights, bias)
    print(f"\nTest Accuracy: {accuracy:.2%}")

    # Plot results
    plot_results(X_train, X_test, y_train, y_test, weights, bias,
        feature_names, (class_a, class_b))

if __name__ == "__main__":
    main()

```

Key Properties

- **Convergence:** Guaranteed for linearly separable data
- **Limitations:** Cannot learn non-linear boundaries
- **Interactive Feature Selection:** Users can choose between sepal/petal features
- **Class Pair Flexibility:** Supports all three binary classification combinations
- **Proper Train-Test Split:** 80-20 ratio with stratified sampling
- **Model Evaluation:** Includes accuracy calculation on test set

4. :for given graph give the following solutions




Figure 1: generalization of intersection

(a) :Generalized Point of Intersection for Shallow Neural Networks for input space parameterized by spherical coordinates θ and ϕ

Solution:

Generalizing the Point of Intersection in Terms of θ and ϕ for Shallow Neural Networks

Step 1: Structure of a Shallow Neural Network

Consider a shallow neural network with:

- Input dimension: d
- Number of **hidden neurons**: m
- Activation function: σ
- Weight vectors: $\mathbf{w}_i \in \mathbb{R}^d$
- Bias terms: $b_i \in \mathbb{R}$
- Output weights: $a_i \in \mathbb{R}$

The output of the network is given by:

$$f(\mathbf{x}) = \sum_{i=1}^m a_i \sigma(\mathbf{w}_i^T \mathbf{x} + b_i)$$

Step 2: Weight Vectors in Angular Coordinates

In spherical coordinates:

$$\mathbf{w} = \|\mathbf{w}\| \begin{bmatrix} \sin(\theta) \cos(\phi) \\ \sin(\theta) \sin(\phi) \\ \cos(\theta) \end{bmatrix}$$

Step 3: Decision Boundary Condition

For each neuron, the decision boundary satisfies:

$$\mathbf{w}_i^T \mathbf{x} + b_i = 0,$$

which in spherical coordinates becomes:

$$\|\mathbf{w}_i\| [x_1 \sin(\theta_i) \cos(\phi_i) + x_2 \sin(\theta_i) \sin(\phi_i) + x_3 \cos(\theta_i)] + b_i = 0$$

Step 4: Intersection of Decision Boundaries

If two neurons intersect, we solve the system:

$$\mathbf{w}_i^T \mathbf{x} + b_i = 0, \quad \mathbf{w}_j^T \mathbf{x} + b_j = 0,$$

which translates to:

$$\|\mathbf{w}_i\| \mathbf{x} \cdot \mathbf{v}(\theta_i, \phi_i) + b_i = 0, \quad \|\mathbf{w}_j\| \mathbf{x} \cdot \mathbf{v}(\theta_j, \phi_j) + b_j = 0$$

Step 5: General Solution

The point of intersection \mathbf{x} can be computed by solving the linear system:

$$\mathbf{x} = \mathbf{A}^{-1} \mathbf{b},$$

where \mathbf{A} is the matrix formed by the weight directions in spherical coordinates, and \mathbf{b} is the bias vector.

- (b) Give the equation of 4 line segments in the graph in terms of θ_1 , θ_2 , θ_3 , etc., for the figure.

Solution:

Consider a shallow neural network with three hidden units and ReLU activations. Let the output y of the network be defined by the following equation:

$$y = \phi_0 + \phi_1 h_1 + \phi_2 h_2 + \phi_3 h_3$$




Figure 2: 4 line equations

where each hidden unit h_i is given by the ReLU activation function:

$$h_i = a(\theta_{i0} + \theta_{i1}x) = \max(0, \theta_{i0} + \theta_{i1}x)$$

The output $y(x)$ is composed of four linear segments, which can be written as:

$$y(x) = \begin{cases} \phi_0, & x < x_1 \\ \phi_0 + \phi_1(\theta_{10} + \theta_{11}x), & x_1 \leq x < x_2 \\ \phi_0 + \phi_1(\theta_{10} + \theta_{11}x) + \phi_2(\theta_{20} + \theta_{21}x), & x_2 \leq x < x_3 \\ \phi_0 + \phi_1(\theta_{10} + \theta_{11}x) + \phi_2(\theta_{20} + \theta_{21}x) + \phi_3(\theta_{30} + \theta_{31}x), & x \geq x_3 \end{cases}$$

Explicitly, the four line segments are:

- First segment: $y = \phi_0$
- Second segment: $y = \phi_0 + \phi_1(\theta_{10} + \theta_{11}x)$
- Third segment: $y = \phi_0 + \phi_1(\theta_{10} + \theta_{11}x) + \phi_2(\theta_{20} + \theta_{21}x)$
- Fourth segment: $y = \phi_0 + \phi_1(\theta_{10} + \theta_{11}x) + \phi_2(\theta_{20} + \theta_{21}x) + \phi_3(\theta_{30} + \theta_{31}x)$

The activation thresholds x_1 , x_2 , and x_3 where each hidden unit is activated are given by:

$$x_i = -\frac{\theta_{i0}}{\theta_{i1}}, \quad \text{for each neuron.}$$

The output function combines the contributions of all active hidden units according to their weights and is expressed in the above piecewise form.

5. : What will be the General Form of the second output in the Two-Output Feedforward Neural Network (2D Case) if one of the output is given ?

Solution:

Two-Output Feedforward Neural Network

We consider a feedforward neural network with:

- 2 input features: x_1, x_2
- D hidden neurons
- 2 output neurons: y_1, y_2
- Activation function $a(\cdot)$ for the hidden layer

1. Hidden Layer Computation

Each hidden unit h_d takes the input vector and applies a linear transformation followed by a nonlinear activation.

$$h_d = a \left(\theta_{d0} + \sum_{i=1}^2 \theta_{di} x_i \right) \quad \text{for } d = 1, 2, \dots, D$$

Where:

- h_d : output of the d -th **hidden unit**
- θ_{d0} : **bias term** for hidden unit d
- θ_{di} : **weight** from input x_i to hidden unit d
- $a(\cdot)$: activation function (e.g., sigmoid, tanh, ReLU)

2. Output Layer Computation

Each output neuron performs a linear combination of all hidden unit outputs with a bias term:

$$y_j = \phi_{j0} + \sum_{d=1}^D \phi_{jd} h_d \quad \text{for } j = 1, 2$$

Where:

- y_j : output of the j -th **output unit**
- ϕ_{j0} : **bias** for output unit j
- ϕ_{jd} : **weight from hidden unit d to output unit j**

3. General equation of both Output

If one of the outputs is:

$$y_1 = \phi_{10} + \sum_{d=1}^D \phi_{1d} h_d$$

Then the other output is similarly given by:

$$y_2 = \phi_{20} + \sum_{d=1}^D \phi_{2d} h_d$$

Explanation:

- Both outputs depend on the same hidden layer outputs h_1, h_2, \dots, h_D
- The only difference lies in the **weights** ϕ_{jd} and **biases** ϕ_{j0} used for each output neuron
- This allows the network to produce different outputs from the same input vector through different weightings

4. Final Combined Form

$$y_1 = \phi_{10} + \sum_{d=1}^D \phi_{1d} \cdot a \left(\theta_{d0} + \sum_{i=1}^2 \theta_{di} x_i \right)$$
$$y_2 = \phi_{20} + \sum_{d=1}^D \phi_{2d} \cdot a \left(\theta_{d0} + \sum_{i=1}^2 \theta_{di} x_i \right)$$

Thus, both outputs are linear combinations of nonlinearly transformed weighted inputs.

6. Let $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$ be independent and identically distributed (i.i.d.) vectors from a multivariate normal distribution:

$$\mathbf{x}_i \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$$

where $\boldsymbol{\mu}$ is the unknown mean vector and Σ is the known covariance matrix.

Solution:

Maximum Likelihood Estimate of Unknown Mean Vector

Let $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$ be independent and identically distributed (i.i.d.) vectors from a multivariate normal distribution:

$$\mathbf{x}_i \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$$

where $\boldsymbol{\mu}$ is the unknown mean vector and Σ is the known covariance matrix.
The probability density function (PDF) of \mathbf{x}_i is given by:

$$f(\mathbf{x}_i|\boldsymbol{\mu}) = \frac{1}{(2\pi)^{p/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x}_i - \boldsymbol{\mu})^T \Sigma^{-1}(\mathbf{x}_i - \boldsymbol{\mu})\right)$$

Likelihood Function

Given the independence of the samples, the likelihood function is the product of the individual densities:

$$L(\boldsymbol{\mu}) = \prod_{i=1}^n f(\mathbf{x}_i|\boldsymbol{\mu})$$

Taking the natural logarithm of the likelihood function (log-likelihood):

$$\log L(\boldsymbol{\mu}) = -\frac{np}{2} \log(2\pi) - \frac{n}{2} \log |\Sigma| - \frac{1}{2} \sum_{i=1}^n (\mathbf{x}_i - \boldsymbol{\mu})^T \Sigma^{-1}(\mathbf{x}_i - \boldsymbol{\mu})$$

Maximizing the Log-Likelihood

To find the MLE of $\boldsymbol{\mu}$, we differentiate $\log L(\boldsymbol{\mu})$ with respect to $\boldsymbol{\mu}$ and set the result to zero:

$$\frac{\partial \log L}{\partial \boldsymbol{\mu}} = \sum_{i=1}^n \Sigma^{-1}(\mathbf{x}_i - \boldsymbol{\mu}) = 0$$

Simplifying:

$$\sum_{i=1}^n (\mathbf{x}_i - \boldsymbol{\mu}) = 0 \implies \sum_{i=1}^n \mathbf{x}_i - n\boldsymbol{\mu} = 0 \implies n\boldsymbol{\mu} = \sum_{i=1}^n \mathbf{x}_i \implies \boldsymbol{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$$

Result: MLE of Mean Vector

$$\hat{\boldsymbol{\mu}} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$$

Thus, the maximum likelihood estimate of the unknown mean vector is the sample mean.

7. The Backpropagation for the cross-entropy loss function of a network of 3 outputs (f_1, f_2, f_3). Just assume that the 3 outputs are the only parameters of the loss function.

Solution:

Cross-Entropy Loss with Softmax Outputs

Let the outputs (logits) of a neural network be f_1, f_2, f_3 . These are passed through the softmax function to produce probabilities:

$$p_i = \frac{e^{f_i}}{\sum_{j=1}^3 e^{f_j}} \quad \text{for } i = 1, 2, 3$$

Let the ground-truth target vector be one-hot encoded: $\mathbf{y} = (y_1, y_2, y_3)$, where $y_k = 1$ for the correct class and $y_i = 0$ for $i \neq k$.

The cross-entropy loss is given by:

$$L = - \sum_{i=1}^3 y_i \log(p_i)$$

Since only one $y_k = 1$, the expression simplifies to:

$$L = -\log(p_k)$$

Computing the Gradient using the Chain Rule

We want to compute the gradient of the loss with respect to the logits f_i i.e. one of the output. Using the chain rule we got:

$$\frac{\partial L}{\partial f_i} = \sum_{j=1}^3 \frac{\partial L}{\partial p_j} \cdot \frac{\partial p_j}{\partial f_i}$$

Step 1: Compute $\frac{\partial L}{\partial p_j}$

$$\frac{\partial L}{\partial p_j} = -\frac{y_j}{p_j}$$

Step 2: Compute $\frac{\partial p_j}{\partial f_i}$ From the derivative of softmax:

$$\frac{\partial p_j}{\partial f_i} = p_j(\delta_{ij} - p_i)$$

Step 3: Combine Using Chain Rule

$$\frac{\partial L}{\partial f_i} = \sum_{j=1}^3 \left(-\frac{y_j}{p_j} \right) \cdot p_j(\delta_{ij} - p_i) = - \sum_{j=1}^3 y_j(\delta_{ij} - p_i)$$

Since $y_j = 1$ only for $j = k$, the summation simplifies:

$$\frac{\partial L}{\partial f_i} = -(\delta_{ik} - p_i) = p_i - y_i$$

Final Result:

$$\frac{\partial L}{\partial f_i} = p_i - y_i \quad \text{for } i = 1, 2, 3$$

Conclusion

This is a well-known result for the gradient of the softmax with cross-entropy loss. It tells us how to adjust each logit during gradient descent.

8. :Backpropagation for 3-class classification using a neural network with 2 inputs, 2 hidden sigmoid units, and 3 softmax output neurons. Derive the forward and backward pass expressions assuming cross-entropy loss.

Solution:

Network Architecture

- **Input layer:** 2 features (x_1, x_2)
- **Hidden layer:** 2 neurons with sigmoid activation
- **Output layer:** 3 neurons with softmax activation (for 3-class classification)
- **Loss function:** Cross-entropy

Notation:

- $W^{[1]} \in \mathbb{R}^{2 \times 2}$: weights from **input to hidden layer**
- $b^{[1]} \in \mathbb{R}^2$: **biases** for hidden layer
- $z^{[1]} \in \mathbb{R}^2$: **pre-activation** of hidden layer
- $a^{[1]} \in \mathbb{R}^2$: activation of hidden layer (**after sigmoid**)
- $W^{[2]} \in \mathbb{R}^{3 \times 2}$: weights **from hidden to output layer**
- $b^{[2]} \in \mathbb{R}^3$: **biases for output layer**
- $z^{[2]} \in \mathbb{R}^3$: **pre-activation** of output layer
- $\hat{y} \in \mathbb{R}^3$: **predicted output probabilities (softmax)**
- $y \in \mathbb{R}^3$: **true label** (one-hot vector)

Forward Pass Equations

1. Hidden Layer Computation

Let:

$$W^{[1]} \in \mathbb{R}^{2 \times 2}, \quad b^{[1]} \in \mathbb{R}^2$$

Then:

$$z^{[1]} = W^{[1]}x + b^{[1]}, \quad a^{[1]} = \sigma(z^{[1]})$$

2. Output Layer Computation

Let:

$$W^{[2]} \in \mathbb{R}^{3 \times 2}, \quad b^{[2]} \in \mathbb{R}^3$$

Then:

$$z^{[2]} = W^{[2]}a^{[1]} + b^{[2]}, \quad \hat{y} = \text{softmax}(z^{[2]})$$

Loss Function

For a true class vector $y \in \mathbb{R}^3$ (one-hot encoded):

$$\mathcal{L} = - \sum_{i=1}^3 y_i \log(\hat{y}_i)$$

Backward Pass (Gradient Computation)

1. Gradient w.r.t. Output Layer (Softmax + Cross Entropy)

$$\frac{\partial \mathcal{L}}{\partial z^{[2]}} = \hat{y} - y$$

2. Gradients for Output Layer Weights and Biases

$$\frac{\partial \mathcal{L}}{\partial W^{[2]}} = (\hat{y} - y) (a^{[1]})^T, \quad \frac{\partial \mathcal{L}}{\partial b^{[2]}} = \hat{y} - y$$

3. Backpropagate to Hidden Layer

$$\delta^{[1]} = \left((W^{[2]})^T (\hat{y} - y) \right) \circ \sigma'(z^{[1]})$$

Where \circ denotes element-wise multiplication and:

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

4. Gradients for Hidden Layer Weights and Biases

$$\frac{\partial \mathcal{L}}{\partial W^{[1]}} = \delta^{[1]} x^T \quad , \quad \frac{\partial \mathcal{L}}{\partial b^{[1]}} = \delta^{[1]}$$

Conclusion

The steps shown above give detail about the backpropagation for a 3-class classification network with 2 inputs, 2 hidden sigmoid units, and a softmax output layer.